

Dimensionality reduction

Machine Learning and Deep Learning

Matteo Boschini, Lorenzo Bonicelli, Aniello Panariello, Emanuele Frascaroli

November 15th, 2022

University of Modena and Reggio Emilia

Principal Component Analysis

Eigenfaces

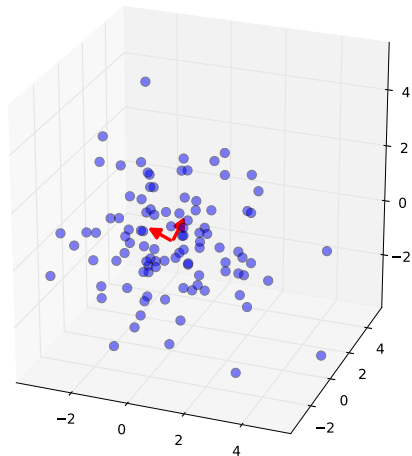
Principal Component Analysis

- Linear dimensionality reduction model
 - Subspace projection is linear
 - Reconstruction is linear
- Projects data in a new space subject to:
 - the direction exhibiting highest variance in feature space is projected on the first axis, the one exhibiting the second highest variance on the second axis, and so on.
 - axis of the new space are orthogonal (covariance is zero).

- Arrange your data in a $n \times d$ matrix X , where n is the number of samples and d is data dimensionality
- Compute the mean μ (d -dimensional vector) of all samples
- Compute covariance matrix


$$\Sigma = (X - \mu)^T (X - \mu)$$

- Pick the first m eigenvectors of Σ (ordered by decreasing eigenvalues), where m is the dimensionality you want your data to be projected to
- Arrange such eigenvectors in a $d \times m$ matrix E
- Compute the projected samples as $P = X \cdot E$
- You can compute the reconstruction as $\tilde{X} = P \cdot E^T$

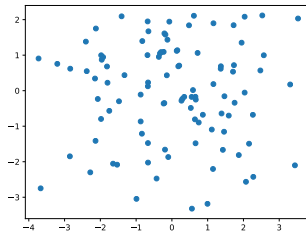


PCA: projecting and reconstructing (2D)


original data



projection

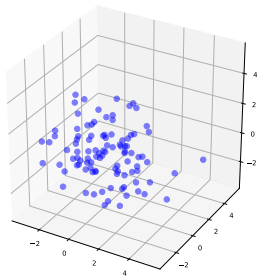


reconstruction

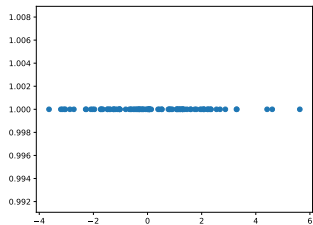


PCA: projecting and reconstructing (1D)

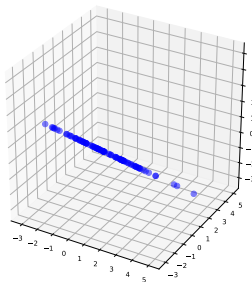
original data



projection



reconstruction



Eigenfaces

Famous algorithm for face recognition. Training is as simple as:

- load faces and annotations from the Olivetti dataset (`datasets.get_faces_dataset` takes care of loading and flattening images)



- Select a number of principal components and fit a PCA on training faces

To classify a test image:

- Project the image in the reduced spaces built in the training phase
- Perform **nearest neighbor classification**:
 - Roughly speaking, choose the class of the nearest training example (in the reduced space)

Each Olivetti image is 112×92 . Once flattened, is a vector of 10304 pixels:

- The covariance matrix is 10304×10304
- Computing eigenvectors and eigenvalues is a pain
- Instead, compute the covariance matrix of transposed X :

$$\Sigma = (X - \mu)(X - \mu)^T$$

- Once selected the principal components \tilde{E} of this weirdo space, you can compute the original eigenvectors just like:


$$E = X^T \cdot \tilde{E}$$

- Normalize the retrieved eigenvectors to have unit length:






$$E_i \leftarrow \frac{E_i}{\sqrt{\lambda_i}} \quad i = 1, 2, \dots, m$$

where λ_i is the eigenvalue corresponding to the eigenvector E_i

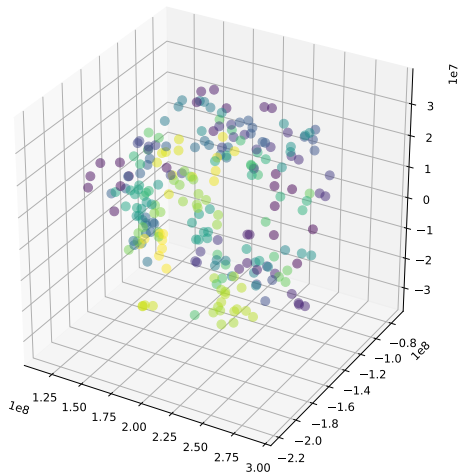
- Mean face:



- Eigenvectors:



Eigenfaces: face space



Eigenfaces: how many dimensions?

